

BigQuery Storage & Spark DataFrames

November 25, 2024

0.1 Scala Version

```
[8]: !scala -version
```

Scala code runner version 2.12.10 -- Copyright 2002-2019, LAMP/EPFL and Lightbend, Inc.

0.2 Creating Spark Session

```
[9]: from pyspark.sql import SparkSession
spark = SparkSession.builder \
    .appName('1.1. BigQuery Storage & Spark DataFrames - Python') \
    .config('spark.jars', 'gs://spark-lib/bigquery/spark-bigquery-latest_2.12.
    ↪jar') \
    .getOrCreate()
```

0.3 Enable repl.eagerEval

```
[10]: # This will output the results of DataFrames in each step without the need
    ↪to show df.show() and also improves the formatting of the output

spark.conf.set("spark.sql.repl.eagerEval.enabled", True)
```

0.4 Reading BigQuery table into Spark DataFrame

```
[11]: # Filtered for english version of Wikipedia for both desktop and mobile versions

table = "bigquery-public-data.wikipedia.pageviews_2020"
df_wiki_pageviews = spark.read \
    .format("bigquery") \
    .option("table", table) \
    .option("filter", "datehour >= '2020-03-01' AND datehour < '2020-03-02'") \
    .load()

df_wiki_pageviews.printSchema()
```

```
root
|-- datehour: timestamp (nullable = true)
```

```

|-- wiki: string (nullable = true)
|-- title: string (nullable = true)
|-- views: long (nullable = true)

```

0.5 Selecting the required columns and filtering for English version

```

[12]: df_wiki_en = df_wiki_pageviews \
      .select("title", "wiki", "views") \
      .where("views > 1000 AND wiki in ('en', 'en.m')") \
      .cache()

```

```
df_wiki_en
```

```

[12]: +-----+-----+-----+
|          title|wiki|  views|
+-----+-----+-----+
|              -| en|143159|
|              -| en| 14969|
|              -| en|186802|
|              -| en|131686|
|              -| en|213787|
|              -| en|211910|
|              -| en|186675|
|              -| en| 21901|
|              -| en|163710|
|              -| en| 23527|
|              -| en|202621|
|              -| en|110524|
|              -| en|220543|
|12_Angry_Men_(195...| en| 1124|
|              -| en|195339|
|              -| en|151283|
|              -| en| 22490|
|              -| en|182985|
|              -| en| 45182|
|              -| en|153327|
+-----+-----+-----+
only showing top 20 rows

```

0.6 Grouping by title and ordering by page views to see the top pages

```

[13]: import pyspark.sql.functions as F

df_wiki_en_totals = df_wiki_en \
  .groupBy("title") \
  .agg(F.sum('views').alias('total_views'))

```

```
df_wiki_en_totals.orderBy('total_views', ascending=False)
```

```
[13]: +-----+-----+
|          title|total_views|
+-----+-----+
|      Main_Page|    10939337|
|United_States_Senate|    5619797|
|              -|    3852360|
|   Special:Search|    1538334|
|2019-20_coronavir...|    407042|
|2020_Democratic_P...|    260093|
|      Coronavirus|    254861|
|The_Invisible_Man...|    233718|
|      Super_Tuesday|    201077|
|      Colin_McRae|    200219|
|      David_Byrne|    189989|
|2019-20_coronavir...|    156803|
|      John_Mulaney|    155605|
|2020_South_Caroli...|    152137|
|      AEW_Revolution|    140503|
|      Boris_Johnson|    120957|
|      Tom_Steyer|    120926|
|Dyatlov_Pass_inci...|    117704|
|      Spanish_flu|    108335|
|2020_coronavirus_...|    107653|
+-----+-----+
only showing top 20 rows
```

0.7 Writing Spark Dataframe to BigQuery table

```
[14]: # Update to your GCS bucket
gcs_bucket = 'andrewmarfo'

# Update to your BigQuery dataset name you created
bq_dataset = 'wikidataset2024'

# Enter BigQuery table name you want to create or overwrite.
# If the table does not exist it will be created when you run the write function
bq_table = 'wiki_total_pageviews'

df_wiki_en_totals.write \
    .format("bigquery") \
    .option("table","{}.{}".format(bq_dataset, bq_table)) \
    .option("temporaryGcsBucket", gcs_bucket) \
    .mode('overwrite') \
    .save()
```

0.8 USING SPARK SQL

0.9 Checking the Scala version

```
[2]: !scala -version
```

Scala code runner version 2.12.10 -- Copyright 2002-2019, LAMP/EPFL and Lightbend, Inc.

0.10 Creating Spark Session

```
[3]: from pyspark.sql import SparkSession
spark = SparkSession.builder \
    .appName('1.2. BigQuery Storage & Spark SQL - Python') \
    .config('spark.jars', 'gs://spark-lib/bigquery/spark-bigquery-latest_2.12.
    ↪jar') \
    .getOrCreate()
```

0.11 Enable repl.eagerEval

```
[4]: # This will output the results of DataFrames in each step without the new need
    ↪to show df.show() and also improves the formatting of the output

spark.conf.set("spark.sql.repl.eagerEval.enabled", True)
```

0.12 Reading BigQuery table into Spark Dataframe

```
[5]: table = "bigquery-public-data.wikipedia.pageviews_2020"
df_wiki_pageviews = spark.read \
    .format("bigquery") \
    .option("table", table) \
    .option("filter", "datehour >= '2020-03-01' AND datehour < '2020-03-02'") \
    .load()

df_wiki_pageviews.printSchema()
```

```
root
 |-- datehour: timestamp (nullable = true)
 |-- wiki: string (nullable = true)
 |-- title: string (nullable = true)
 |-- views: long (nullable = true)
```

0.13 Creating temp table

```
[6]: # Creating temp table to be used in Spark SQL queries
df_wiki_pageviews.createOrReplaceTempView("wiki_pageviews")

# Selecting required columns and applying a filter using WHERE
df_wiki_en = spark.sql("""
SELECT
    title, wiki, views
FROM wiki_pageviews
WHERE views > 1000 AND wiki in ('en', 'en.m')
""").cache()

df_wiki_en
```

```
[6]: +-----+-----+-----+
|          title|wiki| views|
+-----+-----+-----+
|          -| en|143159|
|          -| en| 14969|
|          -| en|186802|
|          -| en|131686|
|          -| en|213787|
|          -| en|211910|
|          -| en|186675|
|          -| en| 21901|
|          -| en|163710|
|          -| en| 23527|
|          -| en|202621|
```

```

|          -| en|110524|
|          -| en|220543|
|12_Angry_Men_(195...| en| 1124|
|          -| en|195339|
|          -| en|151283|
|          -| en| 22490|
|          -| en|182985|
|          -| en| 45182|
|          -| en|153327|
+-----+-----+

```

only showing top 20 rows

0.14 Creating a wiki en pageviews table

```
[7]: df_wiki_en.createOrReplaceTempView("wiki_en")
```

0.15 Grouping by title and finding the top pages by page views

```
[8]: df_wiki_en_totals = spark.sql("""
SELECT
    title,
    SUM(views) as total_views
FROM wiki_en
GROUP BY title
ORDER BY total_views DESC
""")

df_wiki_en_totals
```

```
[8]: +-----+-----+
|          title|total_views|
+-----+-----+
|      Main_Page|    10939337|
|United_States_Senate|    5619797|
|          -|    3852360|
|    Special:Search|    1538334|
|2019-20_coronavir...|    407042|
|2020_Democratic_P...|    260093|
|      Coronavirus|    254861|
|The_Invisible_Man...|    233718|
|      Super_Tuesday|    201077|
|      Colin_McRae|    200219|
|      David_Byrne|    189989|
|2019-20_coronavir...|    156803|
|      John_Mulaney|    155605|
|2020_South_Caroli...|    152137|

```

```
|      AEW_Revolution|      140503|
|      Boris_Johnson|      120957|
|      Tom_Steyer|      120926|
|Dyatlov_Pass_inci...|      117704|
|      Spanish_flu|      108335|
|2020_coronavirus_...|      107653|
+-----+-----+
only showing top 20 rows
```

0.16 Writing Spark Dataframe to BigQuery table

```
[9]: # Updating GCS bucket
gcs_bucket = 'andrewmarfo'

# Updating dataset
bq_dataset = 'wikidataset2024'

# Enter BigQuery table name you want to create or overwrite.
# If the table does not exist it will be created when you run the write function
bq_table = 'wiki_total_pageviews'

df_wiki_en_totals.write \
    .format("bigquery") \
    .option("table", "{}.{}".format(bq_dataset, bq_table)) \
    .option("temporaryGcsBucket", gcs_bucket) \
    .mode('overwrite') \
    .save()
```

0.17 Spark DataFrames & Pandas Plotting - Python

0.17.1 Reading BigQuery table into Spark DataFrame. Filtering to include the date-hour

```
[11]: table = "bigquery-public-data.wikipedia.pageviews_2020"

df_wiki_pageviews = spark.read \
    .format("bigquery") \
    .option("table", table) \
    .option("filter", "datehour >= '2020-03-01' AND datehour < '2020-03-02'") \
    .load()

df_wiki_pageviews.printSchema()
```

```
root
|-- datehour: timestamp (nullable = true)
|-- wiki: string (nullable = true)
|-- title: string (nullable = true)
```

```
|-- views: long (nullable = true)
```

0.18 Selecting required columns and applying a filter using where() which is an alias for filter() then caching the table

```
[12]: df_wiki_en = df_wiki_pageviews \
      .select("datehour", "wiki", "views") \
      .where("views > 1000 AND wiki in ('en', 'en.m')") \
      .cache()

df_wiki_en
```

```
[12]: +-----+-----+-----+
|          datehour|wiki|  views|
+-----+-----+-----+
|2020-03-01 16:00:00|  en|143159|
|2020-03-01 02:00:00|  en| 14969|
|2020-03-01 13:00:00|  en|186802|
|2020-03-01 10:00:00|  en|131686|
|2020-03-01 21:00:00|  en|213787|
|2020-03-01 07:00:00|  en|211910|
|2020-03-01 18:00:00|  en|186675|
|2020-03-01 04:00:00|  en| 21901|
|2020-03-01 15:00:00|  en|163710|
|2020-03-01 01:00:00|  en| 23527|
|2020-03-01 12:00:00|  en|202621|
|2020-03-01 09:00:00|  en|110524|
|2020-03-01 20:00:00|  en|220543|
|2020-03-01 20:00:00|  en|  1124|
|2020-03-01 06:00:00|  en|195339|
|2020-03-01 17:00:00|  en|151283|
|2020-03-01 03:00:00|  en| 22490|
|2020-03-01 14:00:00|  en|182985|
|2020-03-01 00:00:00|  en| 45182|
|2020-03-01 11:00:00|  en|153327|
+-----+-----+-----+
only showing top 20 rows
```

0.19 Grouping by title and ordering by page views to see the top pages

```
[13]: import pyspark.sql.functions as F

df_datehour_totals = df_wiki_en \
  .groupBy("datehour") \
  .agg(F.sum('views').alias('total_views'))
```



```
df_datehour_totals.orderBy('total_views', ascending=False)
```

```
[13]: +-----+-----+
|          datehour|total_views|
+-----+-----+
|2020-03-01 21:00:00|    1642981|
|2020-03-01 06:00:00|    1591160|
|2020-03-01 22:00:00|    1541455|
|2020-03-01 17:00:00|    1535983|
|2020-03-01 18:00:00|    1495387|
|2020-03-01 16:00:00|    1487786|
|2020-03-01 05:00:00|    1469068|
|2020-03-01 07:00:00|    1458756|
|2020-03-01 20:00:00|    1457051|
|2020-03-01 15:00:00|    1446984|
|2020-03-01 19:00:00|    1427811|
|2020-03-01 14:00:00|    1372760|
|2020-03-01 23:00:00|    1353548|
|2020-03-01 08:00:00|    1353292|
|2020-03-01 03:00:00|    1339853|
|2020-03-01 04:00:00|    1312186|
|2020-03-01 12:00:00|    1225647|
|2020-03-01 13:00:00|    1212003|
|2020-03-01 10:00:00|    1211310|
|2020-03-01 09:00:00|    1200977|
+-----+-----+
only showing top 20 rows
```

0.20 Converting Spark DataFrame to Pandas DataFrame

```
[14]: # Converting the Spark DataFrame to Pandas DataFrame and setting the datehour as
      ↳ the index
spark.conf.set("spark.sql.execution.arrow.enabled", "true")
%time pandas_datehour_totals = df_datehour_totals.toPandas()

pandas_datehour_totals.set_index('datehour', inplace=True)
pandas_datehour_totals.head()
```


```
CPU times: user 26.3 ms, sys: 11 ms, total: 37.3 ms
Wall time: 1.89 s
```

```
[14]:          total_views
datehour
2020-03-01 22:00:00    1541455
2020-03-01 09:00:00    1200977
2020-03-01 12:00:00    1225647
2020-03-01 20:00:00    1457051
```

2020-03-01 10:00:00 1211310

0.21 Plotting Pandas Dataframe

```
[16]: import matplotlib.pyplot as plt
pandas_datehour_totals.plot(kind='line',figsize=(12,6));
```



0.22 Plotting Multiple Columns

```
[17]: # Creating a new Spark DataFrame and pivoting the wiki column to create
      ↳ multiple rows for each wiki value
```

```
import pyspark.sql.functions as F

df_wiki_totals = df_wiki_en \
    .groupBy("datehour") \
    .pivot("wiki") \
    .agg(F.sum('views').alias('total_views'))

df_wiki_totals
```

```
[17]: +-----+-----+-----+
|      datehour|    en|  en.m|
+-----+-----+-----+
|2020-03-01 22:00:00|558358|983097|
|2020-03-01 09:00:00|638692|562285|
|2020-03-01 12:00:00|633432|592215|
|2020-03-01 20:00:00|615714|841337|
```

```
|2020-03-01 10:00:00|644680|566630|
|2020-03-01 05:00:00|588808|880260|
|2020-03-01 14:00:00|685500|687260|
|2020-03-01 19:00:00|592967|834844|
|2020-03-01 03:00:00|391300|948553|
|2020-03-01 01:00:00|360511|783510|
|2020-03-01 04:00:00|383489|928697|
|2020-03-01 18:00:00|645590|849797|
|2020-03-01 00:00:00|382154|758920|
|2020-03-01 07:00:00|839531|619225|
|2020-03-01 08:00:00|783419|569873|
|2020-03-01 13:00:00|619111|592892|
|2020-03-01 11:00:00|594027|577016|
|2020-03-01 15:00:00|695881|751103|
|2020-03-01 16:00:00|661878|825908|
|2020-03-01 23:00:00|484077|869471|
+-----+-----+-----+
```

only showing top 20 rows

0.23 Converting to Pandas Dataframe

```
[18]: pandas_wiki_totals = df_wiki_totals.toPandas()

pandas_wiki_totals.set_index('datehour', inplace=True)
pandas_wiki_totals.head()
```


```
[18]:
```

	en	en.m
datehour		
2020-03-01 22:00:00	558358	983097
2020-03-01 09:00:00	638692	562285
2020-03-01 12:00:00	633432	592215
2020-03-01 20:00:00	615714	841337
2020-03-01 10:00:00	644680	566630

0.24 Plotting with line for each column

```
[19]: pandas_wiki_totals.plot(kind='line',figsize=(12,6))
```

```
[19]: <matplotlib.axes._subplots.AxesSubplot at 0x7fa4d0fafa90>
```



[]: